

TP4 - Statistiques sous Python

Analyse de Données Massives - Master 1ère année

Nous allons utiliser dans ce TP le module pandas permettant l'analyse de données avec Python. La première instruction est d'installer le module, à faire dans un terminal de commande. Nous devons aussi installer les modules matplotlib et scipy.

```
{bash}
pip3 install pandas
pip3 install matplotlib
pip3 install scipy
pip3 install numpy
```

Une fois ces modules installée, nous pouvons lancer un notebook pour commencer notre programme.

Il faut tout d'abord importer ces modules. La dernière ligne permettra de voir le résultat des graphiques dans le document.

```
In [160]: import matplotlib.pyplot
import pandas
import scipy.stats
import numpy

%matplotlib inline
```

Données

Nous allons travailler sur les données [tips \(donnees/tips.csv\)](#). Vous pouvez trouver des informations ([ici](#) (<https://www.rdocumentation.org/packages/reshape2/versions/1.4.2/topics/tips>)). Voici comment lire ces données dans python avec `read_csv()` de pandas.

```
In [161]: # Lecture d'un fichier texte
tips = pandas.read_csv("donnees/tips.csv",
header = 0, sep = ",")
```

Sur ces données, il est bien évidemment possible de voir quelques informations classiques.

```
In [162]: type(tips)
Out[162]: pandas.core.frame.DataFrame
```

```
In [163]: # informations diverses  
tips.shape
```

```
Out[163]: (244, 7)
```

```
In [164]: tips.count()
```

```
Out[164]: total_bill    244  
tip          244  
sex          244  
smoker       244  
day          244  
time         244  
size         244  
dtype: int64
```

```
In [165]: tips.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 244 entries, 0 to 243  
Data columns (total 7 columns):  
total_bill    244 non-null float64  
tip          244 non-null float64  
sex          244 non-null object  
smoker       244 non-null object  
day          244 non-null object  
time         244 non-null object  
size         244 non-null int64  
dtypes: float64(2), int64(1), object(4)  
memory usage: 13.4+ KB
```

```
In [166]: list(tips.columns)
```

```
Out[166]: ['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size']
```

```
In [167]: list(tips)
```

```
Out[167]: ['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size']
```

Statistiques descriptives univariés

La fonction `describe()` permet de décrire toutes les variables quantitatives d'un jeu de données directement.

```
In [168]: # résumé basique  
tips.describe()
```

Out[168]:

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

```
In [169]: tips.describe().round(2)
```

Out[169]:

	total_bill	tip	size
count	244.00	244.00	244.00
mean	19.79	3.00	2.57
std	8.90	1.38	0.95
min	3.07	1.00	1.00
25%	13.35	2.00	2.00
50%	17.80	2.90	2.00
75%	24.13	3.56	3.00
max	50.81	10.00	6.00

Quantitative

Il est possible de sélectionner les variables soit via les crochets [], soit par un point ..

Les fonctions ci-dessous permettent de décrire une variable quantitative (ici "total_bill").

```
In [170]: tips.total_bill.describe()
```

```
Out[170]: count    244.000000
          mean     19.785943
          std      8.902412
          min      3.070000
          25%     13.347500
          50%     17.795000
          75%     24.127500
          max     50.810000
          Name: total_bill, dtype: float64
```

```
In [171]: tips["total_bill"].describe()
```

```
Out[171]: count    244.000000
          mean     19.785943
          std      8.902412
          min      3.070000
          25%     13.347500
          50%     17.795000
          75%     24.127500
          max     50.810000
          Name: total_bill, dtype: float64
```

```
In [172]: tips.total_bill.mean()
```

```
Out[172]: 19.78594262295082
```

```
In [173]: tips.total_bill.std()
```

```
Out[173]: 8.9024119548568557
```

```
In [174]: tips.total_bill.var()
```

```
Out[174]: 79.252938613978273
```

```
In [175]: tips.total_bill.min()
```

```
Out[175]: 3.0699999999999998
```

```
In [176]: tips.total_bill.max()
```

```
Out[176]: 50.81000000000002
```

```
In [177]: tips.total_bill.median()
```

```
Out[177]: 17.795
```

```
In [178]: tips.total_bill.quantile([.01, .1, .9, .99])
```

```
Out[178]: 0.01    7.250
          0.10   10.340
          0.90  32.235
          0.99  48.227
Name: total_bill, dtype: float64
```

```
In [179]: scipy.stats.normaltest(tips.total_bill)
```

```
Out[179]: NormaltestResult(statistic=45.117819123473318, pvalue=1.5951078766352608e-1
0)
```

```
In [180]: scipy.stats.shapiro(tips.total_bill)
```

```
Out[180]: (0.9197188019752502, 3.3245434183371003e-10)
```

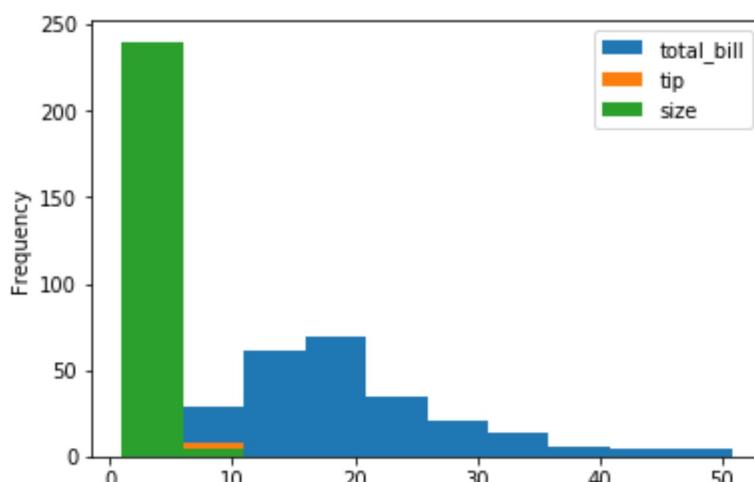
Histogramme

Pour représenter graphiquement cette variable, pandas met à disposition (via le module `matplotlib` utilisé par `pandas`) des fonctions graphiques.

Pour réaliser un **histogramme**, nous utilisons la fonction `hist()`. Celle-ci peut prendre des options. La fonction `plot()` avec le paramètre `kind` avec la valeur "hist" revient au même résultat.

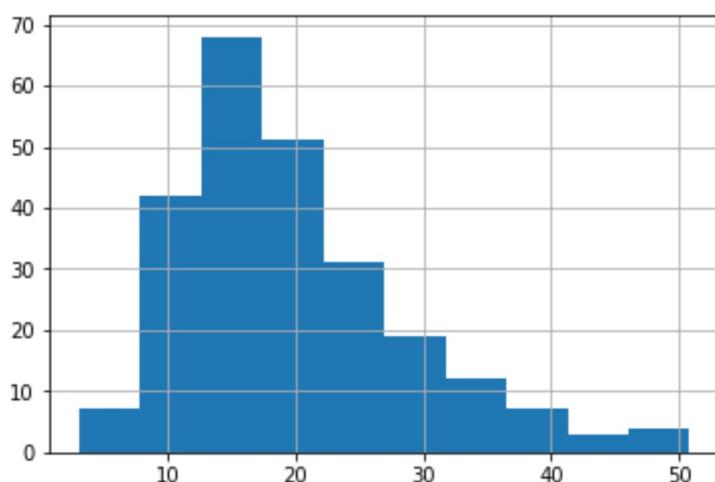
```
In [181]: tips.plot.hist()
```

```
Out[181]: <matplotlib.axes._subplots.AxesSubplot at 0x10d4c4e48>
```



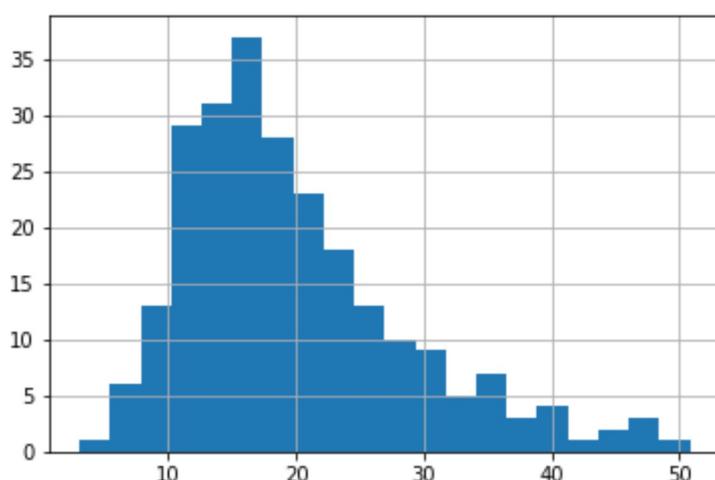
```
In [182]: tips.total_bill.hist()
```

```
Out[182]: <matplotlib.axes._subplots.AxesSubplot at 0x10d4c40b8>
```



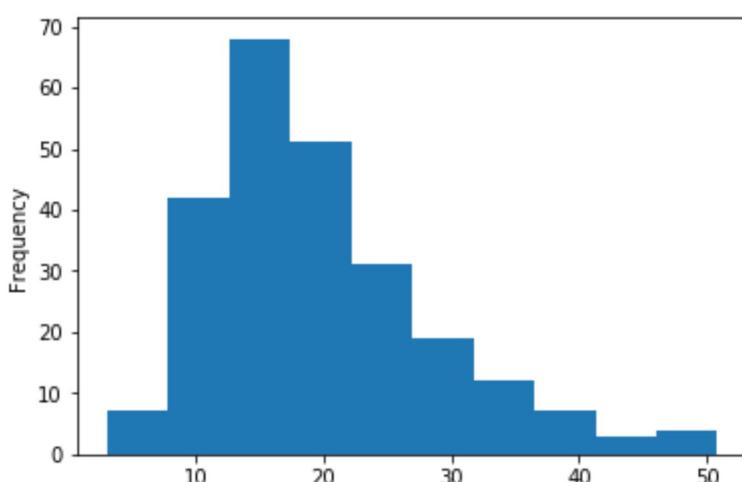
```
In [183]: tips.total_bill.hist(bins = 20)
```

```
Out[183]: <matplotlib.axes._subplots.AxesSubplot at 0x10d687470>
```



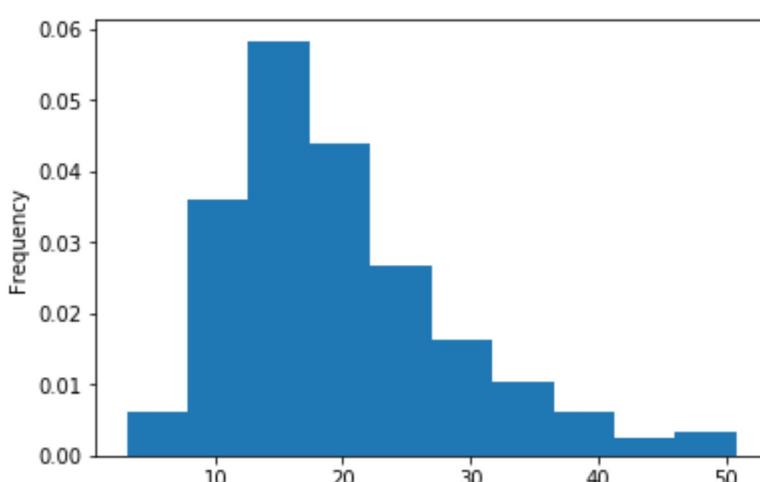
```
In [184]: tips.total_bill.plot(kind = "hist")
```

```
Out[184]: <matplotlib.axes._subplots.AxesSubplot at 0x10d80a860>
```



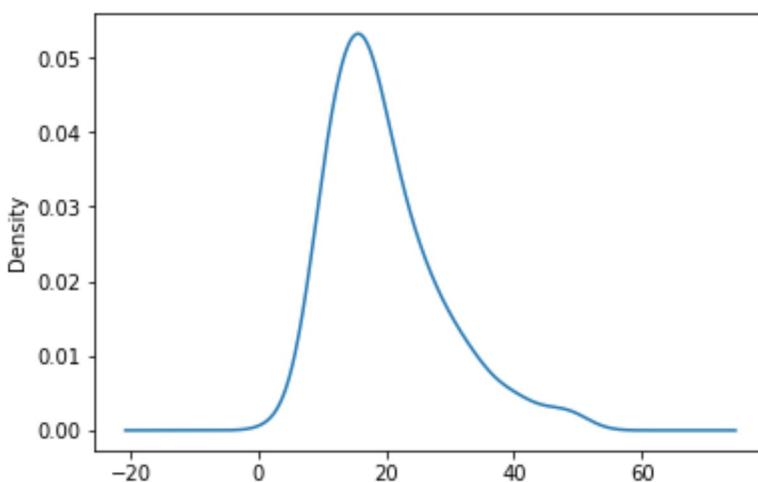
```
In [185]: tips.total_bill.plot(kind = "hist", normed = True)
```

```
Out[185]: <matplotlib.axes._subplots.AxesSubplot at 0x10d90d518>
```



```
In [186]: tips.total_bill.plot(kind = "kde")
```

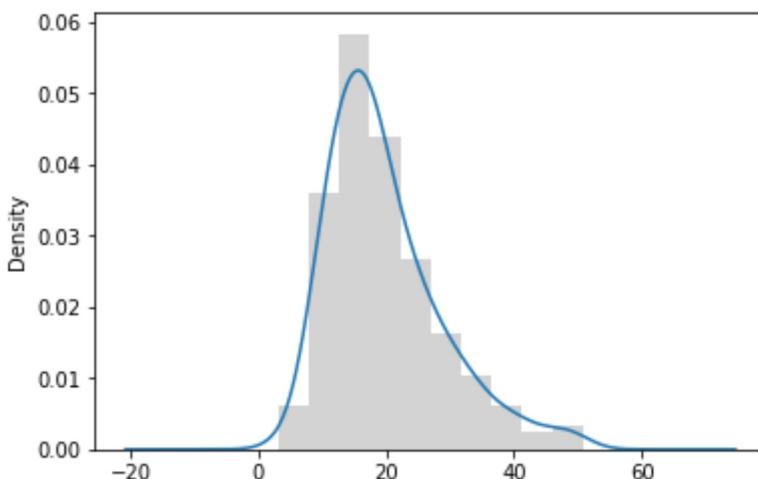
```
Out[186]: <matplotlib.axes._subplots.AxesSubplot at 0x10d9e08d0>
```



Pour avoir la densité et l'histogramme sur le même graphique, il est nécessaire de compiler les deux lignes suivantes **ensemble**.

```
In [187]: tips.total_bill.plot(kind = "hist", normed = True, color = "lightgrey")
tips.total_bill.plot(kind = "kde")
```

```
Out[187]: <matplotlib.axes._subplots.AxesSubplot at 0x10dafb4e0>
```

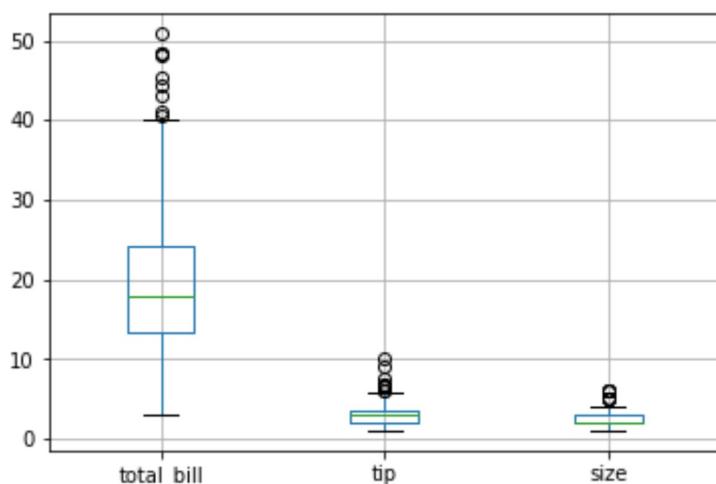


Boîtes à moustaches

Enfin, pour les **boîtes à moustaches**, il faut passer par le DataFrame pour l'afficher, et choisir une variable spécifiquement éventuellement.

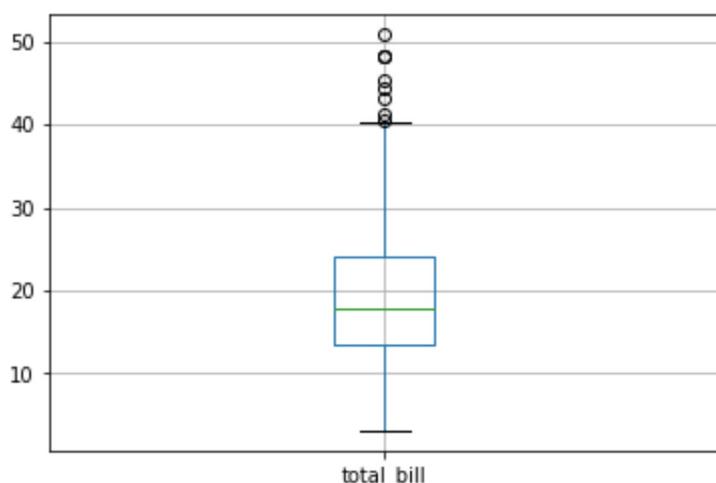
```
In [188]: tips.boxplot()
```

```
Out[188]: <matplotlib.axes._subplots.AxesSubplot at 0x10dc566a0>
```



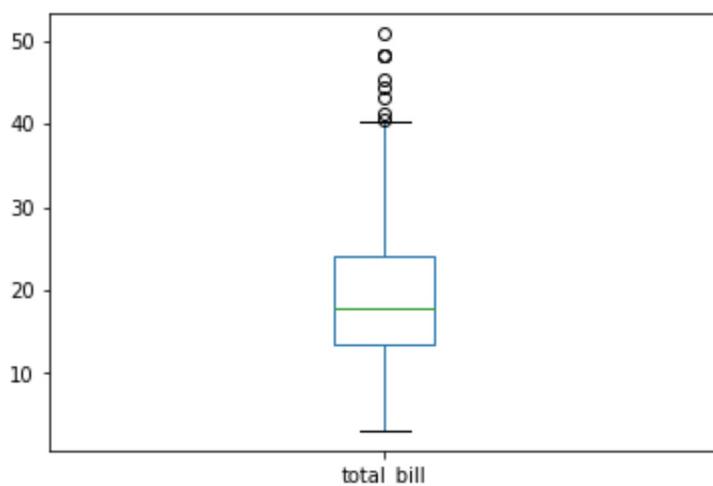
```
In [189]: tips.boxplot(column = "total_bill")
```

```
Out[189]: <matplotlib.axes._subplots.AxesSubplot at 0x10e8c0550>
```



```
In [190]: tips.boxplot(column = "total_bill", grid = False)
```

```
Out[190]: <matplotlib.axes._subplots.AxesSubplot at 0x10e8f1588>
```



Qualitative

Pour les variables qualitatives, il y a plusieurs façons de faire pour obtenir la table d'occurrences (ou des effectifs), ainsi que la table des proportions des modalités.

```
In [191]: tips.sex.describe()
```

```
Out[191]: count      244
unique       2
top        Male
freq       157
Name: sex, dtype: object
```

```
In [192]: tips.sex.unique()
```

```
Out[192]: array(['Female', 'Male'], dtype=object)
```

```
In [193]: tips.sex.value_counts()
```

```
Out[193]: Male      157
Female     87
Name: sex, dtype: int64
```

```
In [194]: pandas.crosstab(tips.sex, "freq")
```

```
Out[194]:
```

col_0	freq
sex	
Female	87
Male	157

```
In [195]: pandas.crosstab(tips.sex, "freq", normalize=True)
```

```
Out[195]:
```

col_0	freq
sex	
Female	0.356557
Male	0.643443

```
In [196]: scipy.stats.chisquare(t)
```

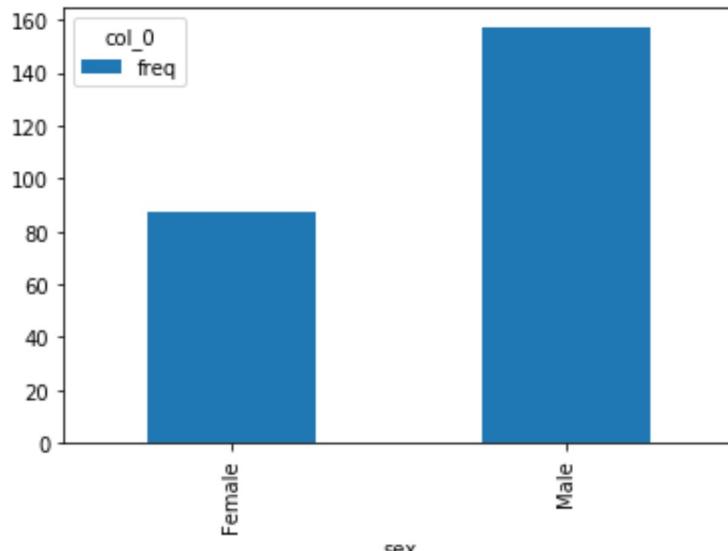
```
Out[196]: Power_divergenceResult(statistic=array([ 20.08196721]), pvalue=array([ 7.41929371e-06]))
```

Diagramme en barres

Ensuite, pour réaliser un **diagramme en barres**, nous utilisons le type "bar" pour `plot()`. Les calculs de proportions précédents nous permettent d'afficher une représentation des proportions plutôt que des effectifs.

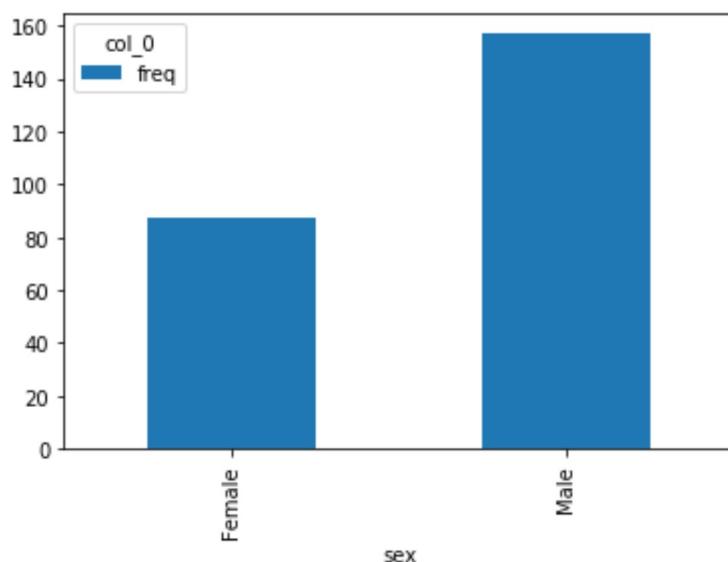
```
In [197]: t = pandas.crosstab(tips.sex, "freq")
t.plot.bar()
```

```
Out[197]: <matplotlib.axes._subplots.AxesSubplot at 0x10e9cd780>
```



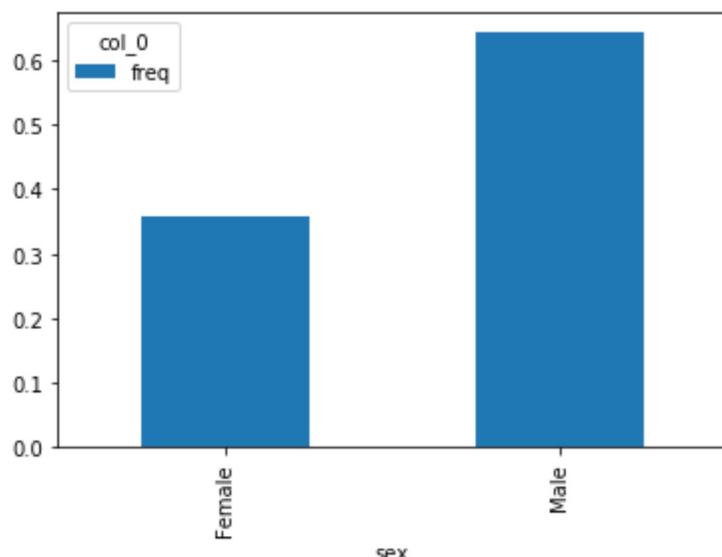
```
In [198]: t.plot(kind = "bar")
```

```
Out[198]: <matplotlib.axes._subplots.AxesSubplot at 0x10ed11908>
```



```
In [199]: t = pandas.crosstab(tips.sex, "freq", normalize=True)  
t.plot(kind = "bar")
```

```
Out[199]: <matplotlib.axes._subplots.AxesSubplot at 0x10ed555f8>
```



```
In [200]: (t * 100).plot(kind = "bar")
```

```
Out[200]: <matplotlib.axes._subplots.AxesSubplot at 0x10e97db00>
```

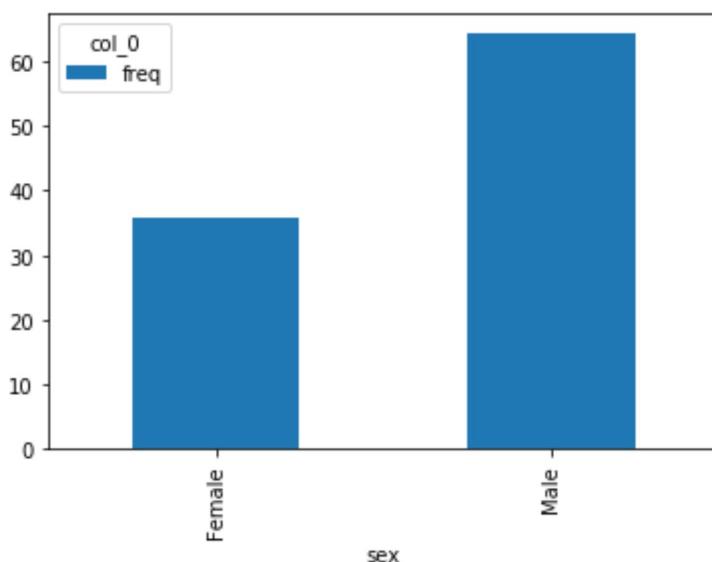
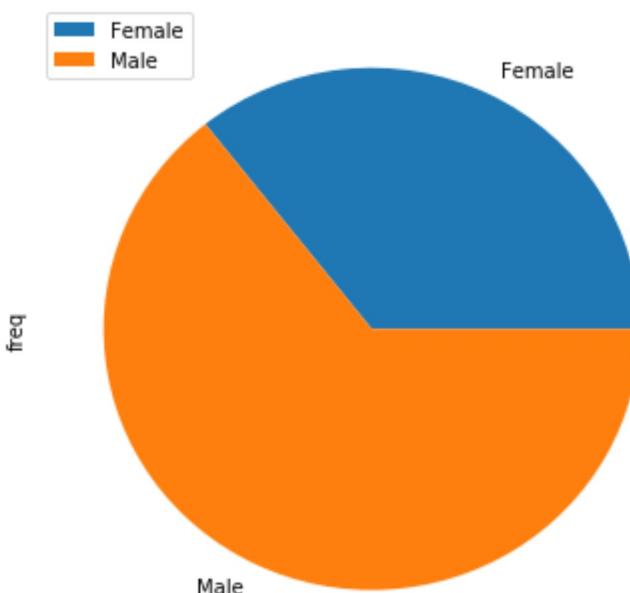


Diagramme circulaire

Et pour un **diagramme circulaire**, seul le tableau des effectifs produit par `value_counts()` nous permet de le réaliser.

```
In [201]: t = pandas.crosstab(tips.sex, "freq")
t.plot.pie(subplots=True, figsize = (6, 6))
```

```
Out[201]: array([<matplotlib.axes._subplots.AxesSubplot object at 0x10efc3470>], dtype=object)
```



Statistiques descriptives bivariées

Quantitative - quantitative

```
In [202]: tips.corr()
```

```
Out[202]:
```

	total_bill	tip	size
total_bill	1.000000	0.675734	0.598315
tip	0.675734	1.000000	0.489299
size	0.598315	0.489299	1.000000

```
In [203]: tips.total_bill.corr(tips.tip)
```

```
Out[203]: 0.67573410921136445
```

```
In [204]: tips.total_bill.cov(tips.tip)
```

```
Out[204]: 8.3235016292248538
```

```
In [205]: scipy.stats.pearsonr(tips.total_bill, tips.tip)
```

```
Out[205]: (0.67573410921136434, 6.6924706468640407e-34)
```

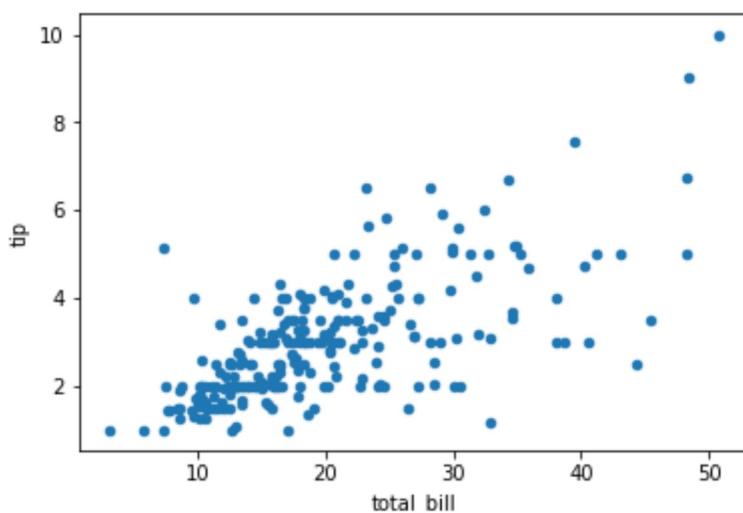
```
In [206]: scipy.stats.kendalltau(tips.total_bill, tips.tip)
```

```
Out[206]: KendalltauResult(correlation=0.51718097214238101, pvalue=2.4455728480214792e-32)
```

Nuage de points

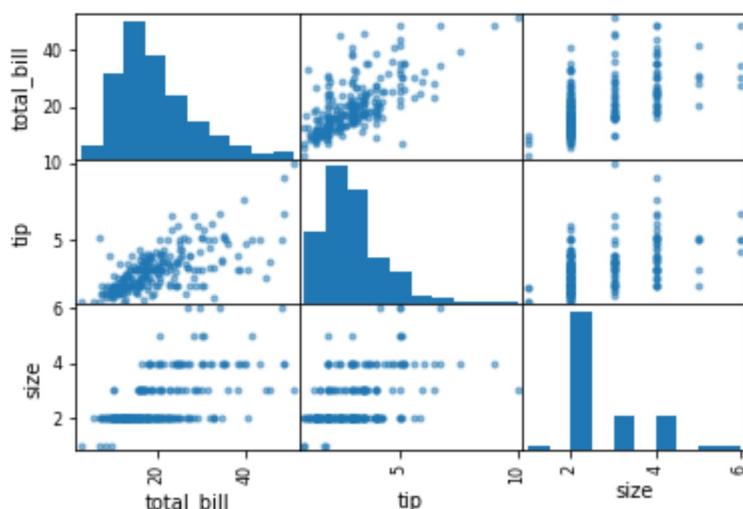
```
In [207]: tips.plot.scatter("total_bill", "tip")
```

```
Out[207]: <matplotlib.axes._subplots.AxesSubplot at 0x10f100e10>
```



```
In [208]: pandas.scatter_matrix(tips)
```

```
Out[208]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x10f23c358>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f388d68>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f4990b8>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x10f4e99b0>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f54fcc0>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f54fcf8>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x10f611860>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f65eb0>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10f6ce128>]], dt
ype=object)
```



Qualitative - qualitative

```
In [209]: pandas.crosstab(tips.sex, tips.smoker)
```

```
Out[209]:
```

smoker	No	Yes
sex		
Female	54	33
Male	97	60

```
In [210]: pandas.crosstab(tips.sex, tips.smoker, margins=True)
```

```
Out[210]:
```

smoker	No	Yes	All
sex			
Female	54	33	87
Male	97	60	157
All	151	93	244

```
In [211]: pandas.crosstab(tips.sex, tips.smoker, normalize = True)
```

```
Out[211]:
```

smoker	No	Yes
sex		
Female	0.221311	0.135246
Male	0.397541	0.245902

```
In [212]: pandas.crosstab(tips.sex, tips.smoker, normalize = "index")
```

```
Out[212]:
```

smoker	No	Yes
sex		
Female	0.620690	0.379310
Male	0.617834	0.382166

```
In [213]: pandas.crosstab(tips.sex, tips.smoker, normalize = "index", margins=True)
```

```
Out[213]:
```

smoker	No	Yes
sex		
Female	0.620690	0.379310
Male	0.617834	0.382166
All	0.618852	0.381148

```
In [214]: pandas.crosstab(tips.sex, tips.smoker, normalize = "columns")
```

```
Out[214]:
```

smoker	No	Yes
sex		
Female	0.357616	0.354839
Male	0.642384	0.645161

```
In [215]: pandas.crosstab(tips.sex, tips.smoker, normalize = "columns", margins=True)
```

```
Out[215]:
```

smoker	No	Yes	All
sex			
Female	0.357616	0.354839	0.356557
Male	0.642384	0.645161	0.643443

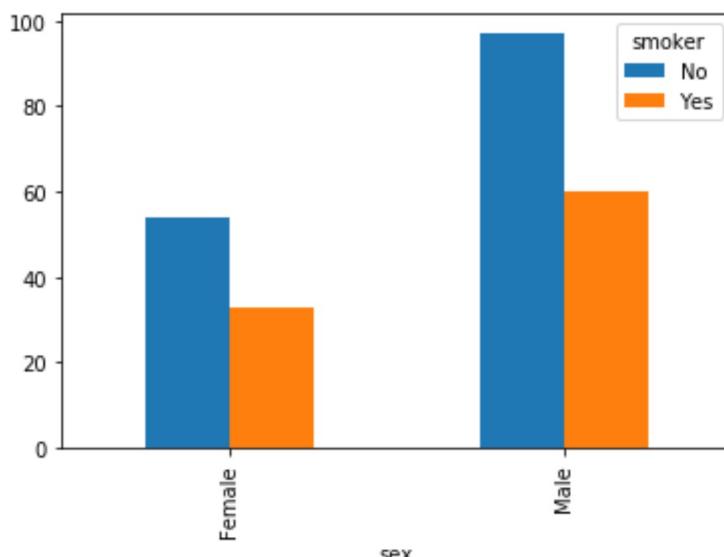
```
In [216]: t = pandas.crosstab(tips.sex, tips.smoker)
scipy.stats.chi2_contingency(t)
```

```
Out[216]: (0.0087632905317735939,
0.92541702049442298,
1,
array([[ 53.84016393,  33.15983607],
       [ 97.15983607,  59.84016393]]))
```

Diagramme en barres

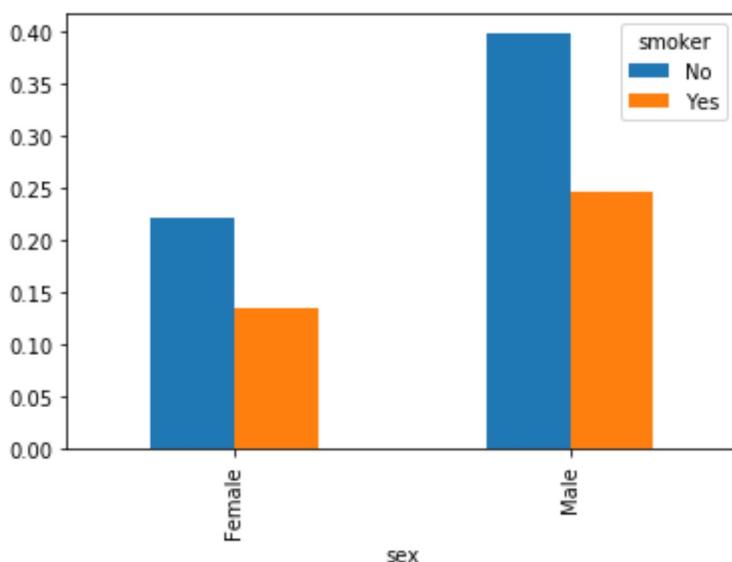
```
In [217]: t = pandas.crosstab(tips.sex, tips.smoker)
t.plot.bar()
```

```
Out[217]: <matplotlib.axes._subplots.AxesSubplot at 0x10f8045f8>
```



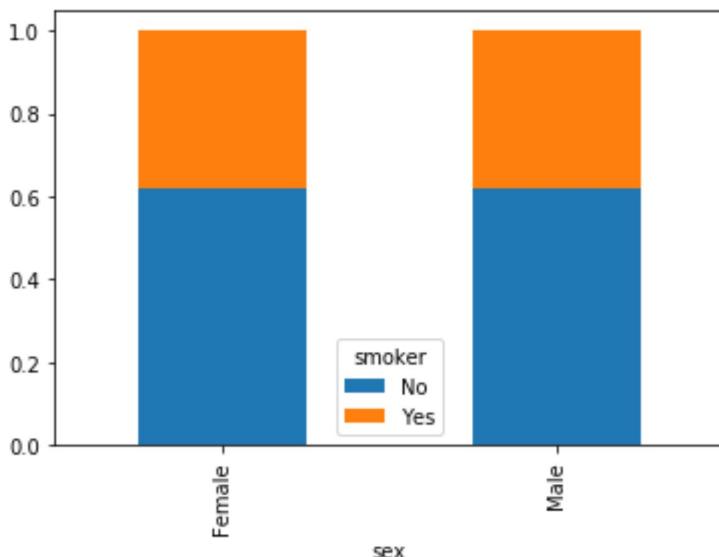
```
In [218]: t = pandas.crosstab(tips.sex, tips.smoker, normalize=True)
t.plot.bar()
```

```
Out[218]: <matplotlib.axes._subplots.AxesSubplot at 0x10f814d68>
```



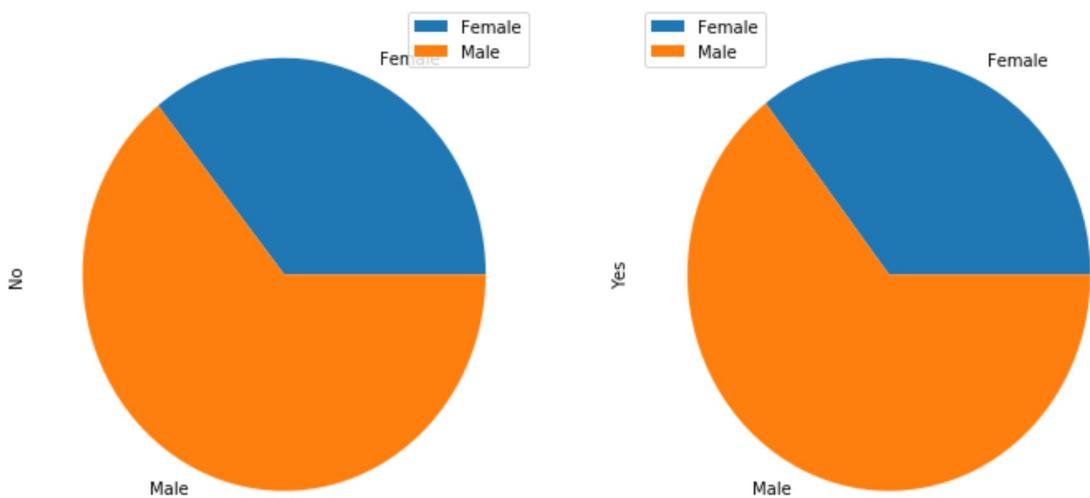
```
In [219]: t = pandas.crosstab(tips.sex, tips.smoker, normalize="index")
t.plot.bar(stacked=True)
```

```
Out[219]: <matplotlib.axes._subplots.AxesSubplot at 0x10faa5a58>
```



```
In [220]: t = pandas.crosstab(tips.sex, tips.smoker)
t.plot.pie(subplots=True, figsize = (12, 6))
```

```
Out[220]: array([<matplotlib.axes._subplots.AxesSubplot object at 0x10faf84a8>,
   <matplotlib.axes._subplots.AxesSubplot object at 0x10fc82d68>], dtype=object)
```



Qualitative - quantitative

```
In [221]: tips.groupby("sex").mean()
```

```
Out[221]:
```

	total_bill	tip	size
sex			
Female	18.056897	2.833448	2.459770
Male	20.744076	3.089618	2.630573

```
In [222]: tips.groupby("sex")["total_bill"].agg([numpy.mean, numpy.std, numpy.median,
numpy.min, numpy.max])
```

```
Out[222]:
```

	mean	std	median	amin	amax
sex					
Female	18.056897	8.009209	16.40	3.07	44.30
Male	20.744076	9.246469	18.35	7.25	50.81

```
In [223]: billFemale = tips.total_bill[tips.sex == "Female"]
billMale = tips.total_bill[tips.sex == "Male"]
scipy.stats.ttest_ind(billFemale, billMale)
```

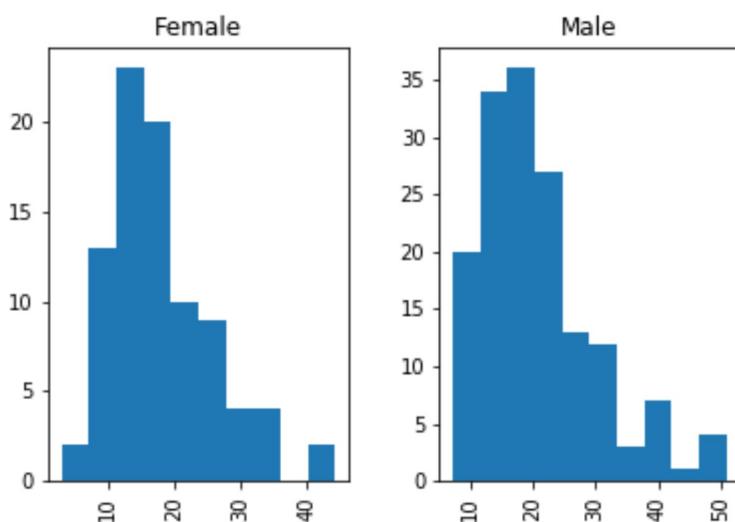
```
Out[223]: Ttest_indResult(statistic=-2.2777940289803134, pvalue=0.023611666846859398)
```

```
In [224]: billGrouped = [tips.total_bill[tips.sex == s] for s in list(tips.sex.unique())]
scipy.stats.f_oneway(*billGrouped)
```

```
Out[224]: F_onewayResult(statistic=5.1883456384583608, pvalue=0.023611666846859697)
```

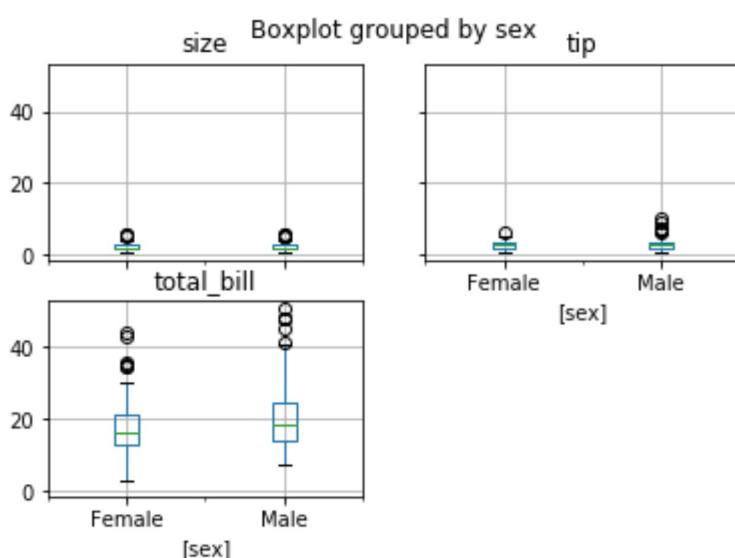
```
In [225]: tips.hist(column = "total_bill", by = "sex")
```

```
Out[225]: array([<matplotlib.axes._subplots.AxesSubplot object at 0x10fdb3cc0>,
                 <matplotlib.axes._subplots.AxesSubplot object at 0x10ff23048>], dtype=object)
```



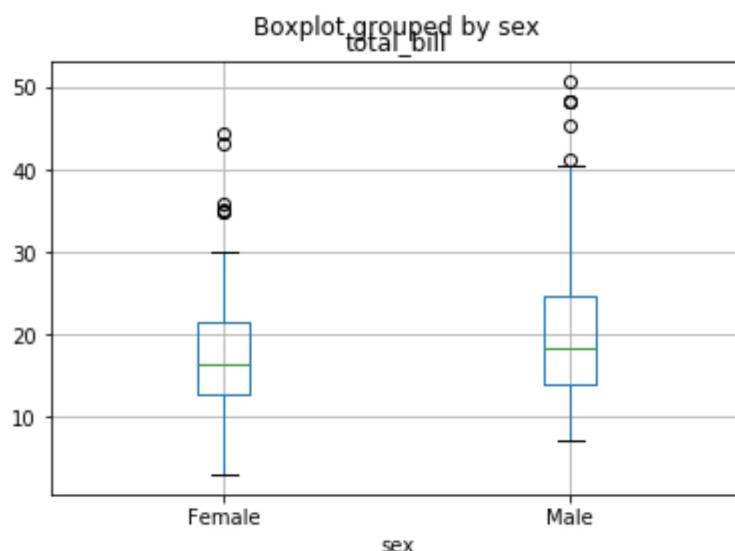
```
In [226]: tips.boxplot(by = "sex")
```

```
Out[226]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x10ffe20b8>,
                  <matplotlib.axes._subplots.AxesSubplot object at 0x10ffffa908>],
                  [<matplotlib.axes._subplots.AxesSubplot object at 0x11017cb70>,
                   <matplotlib.axes._subplots.AxesSubplot object at 0x1101d7208>]], dtype=object)
```



```
In [227]: tips.boxplot(column = "total_bill", by = "sex")
```

```
Out[227]: <matplotlib.axes._subplots.AxesSubplot at 0x1101542b0>
```



Exercices

A partir du fichier [diamonds.csv](#) ([donnees/diamonds.csv](#)) (voir l'aide [ici](#) (<http://docs.ggplot2.org/0.9.3.1/diamonds.html>)), analyser les données suivant le déroulement classique

1. Description de chaque variable
2. Recherche des liens entre le prix (price) et les autres variables

```
In [ ]:
```